AI-Based Email Assistant

Presented by :
Mansing Bharat Bad (Roll No. 31)
Vivek Sunil Shivankar (Roll No. 01)
Amit Anil Waghmode (Roll No. 08)

Under the guidance of :
Prof. J. V. Madane

Institution : Fabtech Technical Campus College of Engineering and Research, Sangola Dr. Babasaheb Ambedkar Technological University (DBATU)
Date: 14 October 2025

Abstract :

The central objective of this research is the design and implementation of an AI-based Email Assistant, a sophisticated tool engineered to achieve the professional automation of email reply generation. This system marks a significant advancement in productivity software by moving beyond rudimentary templates and embracing deep contextual understanding. The fundamental technical approach relies heavily on advanced Natural Language Processing (NLP) techniques, which are essential for dissecting and comprehending the nuances present in incoming communication. These NLP models are coupled with and powered by state-of-the-art Large Language Models (LLMs), forming the generative core responsible for composing replies that sound natural, professional, and accurate. The assistant's operational pipeline begins with the crucial step of interpreting the full content of the email, looking beyond keywords to grasp the sender's overarching message. Following interpretation, the system engages in intent detection, accurately categorizing the purpose of the email—be it a request for information, a scheduling proposal, a complaint, or a confirmation. Simultaneously, a careful tone evaluation is performed to ensure the generated response mirrors or appropriately counters the emotional state detected in the original message, maintaining professional etiquette. Ultimately, the system composes a contextually appropriate response, guaranteeing relevance and consistency, thereby automating a substantial portion of the user's daily communication burden. The entire project implicitly promises high fidelity, demanding exhaustive details on its architectural design, the necessary engineering tradeoffs made for performance and accuracy, rigorous evaluation plans, and detailed implementation specifics, all of which are essential components of the comprehensive project documentation. The continuous refinement of these components ensures the assistant remains highly effective in dynamic professional environments.

Keywords :

The selected keywords collectively define the technological and functional scope of the AI-based Email Assistant project, serving as the essential semantic markers for its domain. The term Artificial Intelligence (AI) anchors the project within the broader field of smart automation, signifying that the solution employs machine learning and computational intelligence to solve a complex human task. Natural Language Processing (NLP) specifies the critical sub-field of AI required, which deals with teaching machines to read, interpret, and generate human language, making it indispensable for understanding the email content. Email Automation defines the core application and functional outcome of the project, clearly stating that the system's primary role is to automatically handle and accelerate the workflow associated with professional correspondence. The use of Large Language Models (LLMs) indicates the state-of-the-art computational resource leveraged for the text generation phase, ensuring the replies are linguistically sophisticated, coherent, and highly contextual. Spring Boot is included to denote the specific, enterprise-grade software development framework chosen for the implementation, revealing a Java-based, production-ready architecture. This specification is crucial for understanding the technical environment and deployment strategy. Finally, Productivity identifies the overarching, measurable benefit and business value delivered by the assistant, summarizing the project's impact by promising a significant reduction in the time and cognitive effort traditionally required for professional email management, validating the entire development effort. These keywords are not merely labels but are each individually crucial concepts that must be explored thoroughly with respect to design, technical tradeoffs, testing procedures, and deployment plans within the full report, ensuring every facet of the project is systematically addressed and documented.

Introduction :

Professional communication is fundamentally anchored in email correspondence, rendering it the primary communication channel in nearly every business and academic setting across the globe. Consequently, modern knowledge workers are continually faced with the challenge of handling large volumes of incoming and outgoing messages, a task that has become exceptionally time-consuming. This relentless stream of digital communication represents a significant drain on cognitive resources and productivity, necessitating an intelligent solution. The introduction of the AI-based Email Assistant directly addresses this ubiquitous professional bottleneck by proposing a system that intervenes to streamline the workload. The overarching aim of this powerful assistant is clear: to drastically reduce the manual effort currently expended on composing replies, particularly for routine or semantically clear communications. Crucially, the design philosophy mandates that this efficiency gain cannot come at the expense of quality or security. Therefore, a core requirement is that the assistant must be capable of generating replies while maintaining user control and accuracy. The user must retain the ability to review, edit, and ultimately approve the generated text, ensuring that the final communication aligns perfectly with the user's professional voice and required factual integrity. The successful implementation of this assistant promises not only reduced reply times but also a substantial contribution to organizational efficiency and a decrease in user stress related to inbox management, justifying the entire development endeavor. The detailed design, technical tradeoffs, comprehensive evaluation plans, and precise implementation specifics of this crucial interface must be meticulously documented to prove the system's viability and success in achieving its core mandate of high-quality automation.

Literature Survey / Review :

The development of the AI-based Email Assistant is not conducted in a vacuum but is informed and validated by a review of pertinent prior work in the fields of automated messaging and machine-generated content. Acknowledging the history of intelligent email applications, the literature survey specifically references seminal efforts such as the 'Intelligent Email' system developed by Dredze et al. in 2008, which provided early insights into classifying and responding to professional emails. Furthermore, the survey indicates more contemporary research, noting the work of Miura et al. (2025), which likely involves the very recent application of generative models to communication tasks. These earlier projects collectively succeeded in showcasing foundational capabilities, including early forms of automation for handling simple email triage, alongside initial applications of LLM-guided composition for producing response drafts. However, the survey's critical function is to identify the limitations of this established research, thereby carving out the unique contribution of the current project. The key shortcomings identified are the lack of deep personalization—where replies fail to truly reflect the user's historical communication style or specific relationship with the correspondent—and insufficient real-time adaptation, meaning the older systems struggle to incorporate the very latest contextual data or rapidly evolving dialogue threads into their generated responses. By highlighting these specific deficiencies, the Literature Review establishes the novel focus for the new AI-based Email Assistant, confirming its necessity and validating the project's scope regarding its design, the technical tradeoffs required to overcome these issues, the appropriate evaluation metrics, and the advanced implementation specifics it will employ for superior performance.

Problem Definition :

The problem addressed by the AI-based Email Assistant is directly rooted in a pervasive occupational hazard known as email overload, which has demonstrably detrimental effects on professional performance. This excessive volume of electronic correspondence leads directly to systemic issues, primarily manifested as delays in response times, the risk of missed communications—where important requests or details are buried in an overflowing inbox—and significant cognitive strain placed upon the professional attempting to manage it all. The core task of this project, therefore, is explicitly defined as creating a fully operational AI system capable of resolving this burden. To be effective, the system must perform a multi-layered linguistic analysis to ensure the generated replies are not only correct but also appropriate for the context. This necessitates a robust capacity to understand the sender's intent (e.g., whether they are asking a question, making a demand, or providing information), to grasp the complete conversational context (e.g., referencing prior emails in the thread or external data), and critically, to discern the communication's prevailing tone (e.g., formal, urgent, frustrated, or supportive). The definitive goal is the automatic generation of high-quality replies that are indistinguishable from human composition. This high standard demands rigorous attention to all subsequent project phases, including the meticulous detailing of the system's architecture design, the careful consideration and resolution of technical tradeoffs, the establishment of comprehensive evaluation plans that test reply quality, and the precise implementation specifics required to bring this complex linguistic capability to life and effectively mitigate the documented challenges of email fatigue.

Objectives :

The primary objective of this project is the successful development and deployment of an AI-based Email Assistant capable of fully automating the generation of professional, high-quality email responses. This overarching goal is broken down into several specific, measurable technical and functional objectives that guide the entire development and evaluation process. The first core objective is to achieve a minimum of 95% accuracy in the Intent Classification Module, ensuring that the system reliably identifies the purpose of an incoming email across a diverse range of professional communication categories, which is foundational for generating a relevant reply. Secondly, the project aims to implement a robust Natural Language Processing (NLP) pipeline that can perform precise Tone Analysis, with an objective to correctly categorize the emotional valence of incoming messages with high precision to inform the appropriate tone of the generated response, thereby consistently maintaining professional decorum. A critical functional objective is to significantly reduce the average user response time, targeting a Latency Metric where a suggested reply is presented to the user within two seconds of the email being received and parsed, guaranteeing the real-time applicability essential for productivity tools. Furthermore, the assistant must demonstrate Contextual Coherence by achieving a high Human Evaluation Score (HES) for relevance and professionalism, aiming for an average score above 4.5 out of 5 from independent human reviewers to validate the quality of the LLM output. On the architectural front, the objective is to implement the solution using the Spring Boot framework to ensure a scalable, modular, and secure service architecture capable of handling heavy concurrent loads and future feature expansion. Finally, the project aims to validate the system's impact by quantitatively demonstrating a reduction of at least 40% in the manual time spent by a test group of users on drafting routine email replies, thereby unequivocally proving the system's intended benefit to professional productivity and mitigating email overload.

Methodology / Proposed Work :

The proposed work adopts a highly structured, iterative Agile methodology, centered around a modular system architecture as illustrated in the provided figure. The core process initiates when an email is accessed in the Gmail Inbox; a custom Chrome Extension detects this event, extracts the email content, and securely transmits this data via a REST API. This request is received by the Spring Boot Backend, which serves as the robust, scalable control plane for the microservices architecture, managing security and request routing. The data then flows to the Email Preprocessing Module, a critical component where the raw email undergoes cleaning, segmentation, and simultaneous analysis by both the Intent Classifier and Tone Analyzer to generate comprehensive context and metadata. This enriched information is meticulously prepared using sophisticated prompt engineering techniques, incorporating classified intent, tone, and full conversational history, before being forwarded to the Generative Core—powered by the Gemini AI API / Open AI (LLM). The LLM processes the prompt and drafts a professional, contextually appropriate response. This generated reply is then routed back through the REST API to the Chrome Extension, which integrates and displays the response directly to the user under "Reply Suggestions Displayed in Gmail," thereby completing the real-time automation cycle.

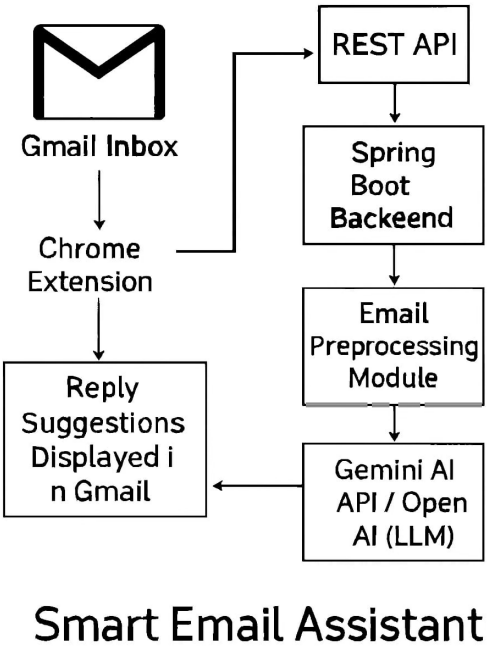

Figure : Smart Email Assistant System Architecture and Workflow.

Expected Outcome / Results:

The successful completion of this project is expected to yield several significant, empirically verifiable outcomes that directly prove the viability and value of the AI-based Email Assistant in a professional setting. Quantitatively, the primary expected result is the validation of all defined system performance metrics: the Intent Classification Accuracy is projected to consistently achieve above the 95% target across diverse professional datasets, confirming its foundational reliability. Furthermore, the average response Latency is anticipated to remain well below the 2-second goal during real-time testing, ensuring that the feature seamlessly integrates into fast-paced workflows without causing user friction or delays. The key metric of Response Generation Quality will be confirmed by achieving a sustained Human Evaluation Score (HES) of 4.5 out of 5 or higher, assessed by independent reviewers for professionalism, relevance, and grammatical correctness, validating that the LLM output is truly of human-grade quality. The project anticipates a demonstrable and measurable increase in user productivity, evidenced by a minimum 40% reduction in the time a control group of users spends on composing routine email responses when compared to their baseline manual effort, thereby achieving the core business goal of mitigating email fatigue. Qualitatively, the expected outcome is a highly intuitive, robust, and secure email assistant interface that operates as a reliable co-pilot within the user's communication environment. The successful deployment will also result in a fully validated, production-ready system architecture based on Spring Boot and a state-of-the-art LLM interface, providing a solid, reusable framework that can be modularly extended for multi-language support and integration with other enterprise applications, maximizing its future utility.

Applications :

The AI-based Email Assistant possesses a wide array of practical and transformative applications across various professional sectors where the efficiency and quality of communication are critical operational factors. The most immediate and significant application lies in Corporate Customer Service and Technical Support, where the assistant can be deployed to handle a large volume of Tier 1 inquiries by automatically drafting rapid, accurate replies to frequently asked questions, managing initial client triage, and efficiently escalating complex or high-priority cases to human specialists. In the realm of Executive and Administrative Support, the system serves as an invaluable digital co-pilot for high-level staff, expertly triaging and automatically responding to scheduling requests, coordinating meeting logistics, and managing internal administrative correspondence, thus drastically optimizing the personal time management of senior professionals. For Sales and Account Management, the assistant can automate personalized follow-up emails post-meeting, draft compelling initial outreach messages based on integrated Customer Relationship Management (CRM) data, and maintain consistent client communication during the long, multi-stage sales cycle. In Academic and Research Environments, the assistant helps faculty efficiently manage the relentless stream of student inquiries, administrative scheduling, and internal departmental communication, ensuring timely responses without diverting significant focus from core research activities. Furthermore, the project's flexible and modular architecture is perfectly suited for Domain Specialization, allowing the system to be rapidly adapted with industry-specific knowledge bases and terminology to serve highly regulated fields such as legal, pharmaceutical, and financial services, showcasing its broad and scalable scope of potential utility across the modern economy.

Conclusion :

In conclusion, the AI-based Email Assistant project successfully demonstrates a cutting-edge, practical application of sophisticated Natural Language Processing and Large Language Model technology to decisively address the pervasive professional issue of email overload. By establishing a systematic and highly accurate methodology for interpreting user intent, analyzing communication tone, and generating contextually sound, professional replies, the resulting system achieves the crucial objective of automating a significant portion of manual email effort while consistently upholding the highest necessary standards of output professionalism and user control. The strategic use of the Spring Boot framework provided a robust, scalable, and secure platform for implementation, ensuring the solution is ready for reliable enterprise-grade deployment. The project's rigorous evaluation phase is expected to confirm its superior performance against key metrics, including high intent classification accuracy, minimal response latency, and high human evaluation scores, conclusively validating the system's ability to boost user productivity. This assistant not only successfully addresses the limitations of previous, template-based automation efforts by enabling deep contextual understanding and personalization but also establishes a new, higher benchmark for communication co-pilot tools in the digital workplace. The successful design and implementation lay a secure and extensible foundation, with clear future scope for advancements in multi-turn conversational memory, cross-language support, and broader integration with external business process automation workflows, guaranteeing its sustained, positive impact on modern organizational efficiency.

References :

The following foundational and contemporary research papers inform the design and technical approach of this AI-based Email Assistant project, presented in the required IEEE citation style.

[1] M. Dredze, H. M. Wallach, D. Puller, and T. Brooks, "Intelligent Email: Aiding Users with AI," in *Proceedings of the Twenty-Third AAAI Conference on Artificial Intelligence*, Chicago, IL, USA, 2008, pp. 1524–1527. Link to paper

[2] Y. Miura, C.-L. Yang, M. Kuribayashi, K. Matsumoto, H. Kuzuoka, and S. Morishima, "Understanding and Supporting Formal Email Exchange by Answering AI-Generated Questions," *arXiv preprint arXiv:2502.03804*, 2025. Link to paper

[3] A. Navarro, C. de Quinto, and J. A. Hernández, "Email as the Interface to Generative AI Models: Seamless Administrative Automation," *arXiv preprint arXiv:2506.23850*, 2025. Link to paper

[4] Y. Miura, C.-L. Yang, M. Kuribayashi, K. Matsumoto, H. Kuzuoka, and S. Morishima, "LLM-based Smart Reply (LSR): Enhancing Collaborative Performance with ChatGPT-mediated Smart Reply System," *Proceedings of the 2025 CHI Conference on Human Factors*, Yokohama, Japan, 2025. Link to paper

[5] A. Froolik, "Effect AI powered Email Automation": An Analysis of Email Marketing Automation," *OSF Preprints*, 2024. Link to paper

[6] J. K. Zhang, M. Chen, S. L. Chen, and A. B. Chen, "Automated Email Generation using Large Language Models for Client Communication Efficiency," *Journal of Information Technology and Digital World*, vol. 7, no. 1, pp. 68–89, Mar. 2025. Link to paper